

## Overview

Your objective is to build a **small AI-driven application** showcasing key elements of **LangChain v0.3** and **LangGraph v0.3**:

1. **Document ingestion and retrieval** (with embeddings and a local vector store).
2. A **multi-step or agentic flow** orchestrated by **LangGraph**.
3. **Memory** for multi-turn context.
4. **Streaming** or partial/step-by-step output.
5. **Basic reliability/error-handling** to demonstrate real-world readiness.

**Note:** You can **choose** any LLM or additional external APIs and tools, as needed. If your solution depends on third-party services (e.g., external data or special APIs), just provide **clear instructions** (e.g., how to obtain credentials or keys). There is **no requirement** that the chosen LLM or external services be free to use, only that your instructions make it feasible for a reviewer to replicate or test your solution if they also have the same LLM account (e.g., an OpenAI API key).

---

## High-Level Requirements

### 1. Document Ingestion & Retrieval (RAG)

- Ingest a small local dataset (a handful of text/Markdown files).
- Split these documents and create embeddings.
- Store them in a local vector store (e.g., ChromaDB) so it's easy to spin up and test.
- Illustrate how you retrieve relevant chunks to ground the LLM's outputs.

### 2. LangChain Components

- Implement at least one **Prompt Template** or a chain using Runnables.
- Show usage of a **tool** (function calling, or any other accessible external function) if you aim for agentic behavior.

### 3. LangGraph Flow

- Define a **graph**.
- **If you claim it's an agent** (like ReAct/CodeAct), ensure it genuinely demonstrates "Think → Act → Observe" or a comparable agentic loop.
- Otherwise, a standard multi-step workflow is also fine, so long as your graph includes some meaningful branching or multi-step transitions.

### 4. Memory

- Maintain conversation or contextual state across multiple user turns.
- This can be done with either the state or via **LangGraph's** more advanced memory/persistence.

### 5. Streaming

- Demonstrate partial or incremental output – e.g., token-by-token streaming from the LLM or real-time updates about the agent's steps.

## 6. Reliability / Error Handling

- Show a small mechanism to handle or log errors (invalid LLM output, timeouts, etc.).
- Optionally, use a callback or a mini fallback node if the system fails to produce expected results.

## 7. User Interface

- Provide a **simple** but testable UI. You can pick **any** approach (CLI, Streamlit, Gradio, minimal web server, etc.) so reviewers can quickly see the system in action.
- Let the user run multiple queries that demonstrate memory across sessions, retrieval from your vector DB, and your multi-step logic.

## 8. Documentation & Setup

- **README** with steps to install, run ingestion (if separate), and start the interface.
- List any environment variables or secrets needed (e.g., an `.env.example`).
- Show how a reviewer can test it (e.g., “Run `python main.py` and ask about X. You’ll see retrieval from the local docs.”).

---

## Deliverables

1. **Source Code** in a well-structured format (repo or archive).
2. **README** detailing installation, usage, plus any setup for external APIs if applicable.
3. **Scripts or Functions** for data ingestion and app execution.
4. **Optional Tests** if you have time (simple integration or unit tests).

---

## Review Criteria

- **Correct Usage of LangChain & LangGraph:** Appropriate use of models, memory, Runnables, compiled graphs, etc.
- **RAG Implementation:** Are you effectively retrieving relevant data from the local DB to ground LLM answers?
- **Agentic / Multi-Step Flow:** If you present an agent, does it genuinely show agentic steps? If not an agent, is your multi-step logic structured, branching, or otherwise meaningful beyond a single LLM call?
- **Memory & Streaming:** Are conversation details retained, and is partial output shown?
- **Reliability:** Basic logging, callbacks, or error handling in place.
- **Documentation:** Clear instructions so others can replicate or test with their own LLM keys or service credentials.

Feel free to add extra flair or features. The goal is to demonstrate a **small but realistic** LLM-based application using the latest **LangChain** and **LangGraph** capabilities in a way that’s straightforward for others to run and review. Good luck!